



## Homework # 11: Support Vector Machines (Chap. 9)

1. (**Chap. 9, # 1a, p. 368**) This problem involves a hyperplane in two dimensions. Sketch the hyperplane  $1 + 3X_1 - X_2 = 0$ . Indicate the set of points for which  $1 + 3X_1 - X_2 > 0$ , as well as the set of points for which  $1 + 3X_1 - X_2 < 0$ .




For the points which  $1 + 3X_1 - X_2 < 0$ , they are colored red. For the points which  $1 + 3X_1 - X_2 > 0$ , they are colored blue.

2. (**Chap. 9, # 2, p. 368**) We have seen that in  $p = 2$  dimensions, a linear decision boundary takes the form  $\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$ . We now investigate a **non-linear** decision boundary.
- (a) Sketch the curve  $(1 + X_1)^2 + (2 - X_2)^2 = 4$ .




- (b) On your sketch, indicate the set of points for which  $(1 + X_1)^2 + (2 - X_2)^2 > 4$ , as well as the set of points for which  $(1 + X_1)^2 + (2 - X_2)^2 \leq 4$ .



Blue points indicate  $(1 + X_1)^2 + (2 - X_2)^2 > 4$ , while red points indicate  $(1 + X_1)^2 + (2 - X_2)^2 \leq 4$ .

- (c) Suppose that a classifier assigns an observation to the blue class if  $(1 + X_1)^2 + (2 - X_2)^2 > 4$ , and to the red class otherwise. To what class is the observation  $(0, 0)$  classified?  $(-1, 1)$ ?  $(2, 2)$ ?  $(3, 8)$ ?



The observation  $(0, 0)$ ,  $(2, 2)$ ,  $(3, 8)$  are classified as blue while  $(-1, 1)$  is classified as red.

- (d) Argue that while the decision boundary in (c) is not linear in terms of  $X_1$  and  $X_2$ , it is linear in terms of  $X_1$ ,  $X_1^2$ ,  $X_2$ , and  $X_2^2$ .


$$\begin{aligned} & (1 + X_1)^2 + (2 - X_2)^2 \\ &= 1 + 2X_1 + X_1^2 + 4 - 4X_2 + X_2^2 \\ &= 5 + 2X_1 + X_1^2 + 4 - 4X_2 + X_2^2 \end{aligned}$$

3. (**Chap. 9, # 3, p. 368**) Here we explore the maximal margin classifier on a toy data set.


- (a) We are given  $n = 7$  observations in  $p = 2$  dimensions. For each observation, there is an associated class label.

Obs.	$X_1$	$X_2$	Y
1	3	4	Red
2	2	2	Red
3	4	4	Red
4	1	4	Red
5	2	1	Blue
6	4	3	Blue
7	4	1	Blue

Sketch the observations.



- (b) Sketch the optimal separating hyperplane, and provide the equation for this hyperplane such as in exercise #1.




- (c) Describe the classification rule for the maximal margin classifier. It should be something along the lines of "Classify to Red if  $\beta_0 + \beta_1 X_1 + \beta_2 X_2 > 0$ , and classify to Blue otherwise". Provide the values for  $\beta_0$ ,  $\beta_1$ , and  $\beta_2$ .

The observation is classified to Red if  $0.5 - X_1 + X_2 > 0$ , and classified to Blue otherwise.  $\beta_0 = 0.5$ ,  $\beta_1 = -1$ , and  $\beta_2 = 1$ .

- (d) On your sketch, indicate the margin for the maximal margin hyperplane.


Maximal margin:  $M = \frac{\sqrt{0.5^2 + 0.5^2}}{2} = \frac{\sqrt{2}}{4}$ .

- (e) Indicate the support vectors for the maximal margin classifier.




The points (2,1), (2,2), (4,3) and (4,4) are the support vectors.

- (f) Argue that a slight movement of the seventh observation would not affect the maximal margin hyperplane.



For the seventh observation, (4,1), a slight movement would not affect the maximal margin hyperplane as long as it does not enter the existed hyperplane.


- (g) Sketch a hyperplane that is not the optimal separating hyperplane, and provide the equation for this hyperplane.



This hyperplane is not optimal since it does not generate the largest margin, and the equation for this hyperplane is  $0.3 - 0.9X_1 + X_2 = 0$

- (h) Draw an additional observation on the plot so that the two classes are no longer separable by a hyperplane.





4. (Chap. 9, # 7, p. 371) In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the Auto data set (from ISLR).

- i. Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.

```
Auto$mileage <- rep(1,nrow(Auto))
Auto$mileage[Auto$mpg < median(Auto$mpg)] <- 0
Auto$mileage <- as.factor(Auto$mileage)
```

- ii. Fit a support vector classifier to the data with various values of cost, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results.

```
set.seed(444)
s.b <- tune(svm, mileage ~ weight + year, data = Auto,
           ranges = list(cost = 10^seq(-3,3,1), kernel = 'linear'))
summary(s.b)

##
```

```
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost kernel
##   0.1 linear
##
## - best performance: 0.09192308
##
## - Detailed performance results:
##   cost kernel      error dispersion
## 1 1e-03 linear 0.41846154 0.09093521
## 2 1e-02 linear 0.10461538 0.05974185
## 3 1e-01 linear 0.09192308 0.04399032
## 4 1e+00 linear 0.09955128 0.05461402
## 5 1e+01 linear 0.09698718 0.04946574
## 6 1e+02 linear 0.09955128 0.05461402
## 7 1e+03 linear 0.09955128 0.05461402

s.b.2 <- tune(svm, mileage ~ weight + year, data = Auto,
              ranges = list(cost = seq(0.01,0.2,0.005), kernel = 'linear'))
summary(s.b.2)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost kernel
##   0.02 linear
##
## - best performance: 0.08935897
##
## - Detailed performance results:
##   cost kernel      error dispersion
## 1 0.010 linear 0.10205128 0.07645441
## 2 0.015 linear 0.09448718 0.08027832
## 3 0.020 linear 0.08935897 0.05832057
## 4 0.025 linear 0.09448718 0.05420465
## 5 0.030 linear 0.09705128 0.05388701
## 6 0.035 linear 0.09455128 0.05431324
## 7 0.040 linear 0.09455128 0.05431324
## 8 0.045 linear 0.09455128 0.05431324
```

```
## 9 0.050 linear 0.09711538 0.05399286
## 10 0.055 linear 0.09967949 0.05488189
## 11 0.060 linear 0.09967949 0.05488189
## 12 0.065 linear 0.09967949 0.05619719
## 13 0.070 linear 0.09711538 0.05399286
## 14 0.075 linear 0.09711538 0.05399286
## 15 0.080 linear 0.09967949 0.05619719
## 16 0.085 linear 0.09711538 0.05399286
## 17 0.090 linear 0.09711538 0.05399286
## 18 0.095 linear 0.09711538 0.05399286
## 19 0.100 linear 0.09711538 0.05399286
## 20 0.105 linear 0.09455128 0.05431324
## 21 0.110 linear 0.09711538 0.05399286
## 22 0.115 linear 0.09711538 0.05399286
## 23 0.120 linear 0.09711538 0.05399286
## 24 0.125 linear 0.09711538 0.05399286
## 25 0.130 linear 0.09711538 0.05399286
## 26 0.135 linear 0.09711538 0.05399286
## 27 0.140 linear 0.09711538 0.05399286
## 28 0.145 linear 0.09711538 0.05399286
## 29 0.150 linear 0.09711538 0.05399286
## 30 0.155 linear 0.09967949 0.05619719
## 31 0.160 linear 0.09967949 0.05619719
## 32 0.165 linear 0.09711538 0.05663421
## 33 0.170 linear 0.09455128 0.05431324
## 34 0.175 linear 0.09455128 0.05431324
## 35 0.180 linear 0.09455128 0.05431324
## 36 0.185 linear 0.09455128 0.05431324
## 37 0.190 linear 0.09455128 0.05431324
## 38 0.195 linear 0.09711538 0.05399286
## 39 0.200 linear 0.09455128 0.05431324
```

The best cost from the output of the 10-fold cross validation is 0.02, with the lowest cross-validation error 0.08935897.

- iii. Now repeat (b), this time using SVMs with radial and polynomial basis kernels. Comment on your results.

```
set.seed(444)
s.c <- tune(svm, mileage ~ weight + year, data = Auto, ranges = list(
  cost = seq(0.01,0.2,0.005), kernel = c('linear','radial','polynomial'))
summary(s.c)

##
## Parameter tuning of 'svm':
##
```

```
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost kernel
##   0.025 linear
##
## - best performance: 0.08423077
##
## - Detailed performance results:
##   cost      kernel      error dispersion
## 1    0.010    linear 0.10461538 0.05974185
## 2    0.015    linear 0.09705128 0.05388701
## 3    0.020    linear 0.09192308 0.05019521
## 4    0.025    linear 0.08423077 0.04689205
## 5    0.030    linear 0.08679487 0.04867313
## 6    0.035    linear 0.08935897 0.04724900
## 7    0.040    linear 0.09192308 0.04871814
## 8    0.045    linear 0.08935897 0.04724900
## 9    0.050    linear 0.08935897 0.04724900
## 10   0.055    linear 0.08935897 0.04724900
## 11   0.060    linear 0.09192308 0.05163005
## 12   0.065    linear 0.09192308 0.05163005
## 13   0.070    linear 0.09192308 0.04399032
## 14   0.075    linear 0.09192308 0.04399032
## 15   0.080    linear 0.09192308 0.04399032
## 16   0.085    linear 0.09192308 0.04399032
## 17   0.090    linear 0.09192308 0.04399032
## 18   0.095    linear 0.09192308 0.04399032
## 19   0.100    linear 0.09192308 0.04399032
## 20   0.105    linear 0.09192308 0.04399032
## 21   0.110    linear 0.09192308 0.04399032
## 22   0.115    linear 0.09192308 0.04399032
## 23   0.120    linear 0.09192308 0.04399032
## 24   0.125    linear 0.09192308 0.04399032
## 25   0.130    linear 0.09192308 0.04399032
## 26   0.135    linear 0.09192308 0.04399032
## 27   0.140    linear 0.09192308 0.04399032
## 28   0.145    linear 0.09442308 0.04519425
## 29   0.150    linear 0.09442308 0.04519425
## 30   0.155    linear 0.09192308 0.04399032
## 31   0.160    linear 0.09192308 0.04399032
## 32   0.165    linear 0.09192308 0.04399032
## 33   0.170    linear 0.09192308 0.04399032
## 34   0.175    linear 0.09192308 0.04399032
```

##	35	0.180	linear	0.09192308	0.04399032
##	36	0.185	linear	0.09192308	0.04399032
##	37	0.190	linear	0.09192308	0.04399032
##	38	0.195	linear	0.09192308	0.04399032
##	39	0.200	linear	0.09192308	0.04399032
##	40	0.010	radial	0.14051282	0.06110723
##	41	0.015	radial	0.09955128	0.05984588
##	42	0.020	radial	0.09442308	0.05679586
##	43	0.025	radial	0.08935897	0.05174930
##	44	0.030	radial	0.08935897	0.05174930
##	45	0.035	radial	0.09185897	0.05158214
##	46	0.040	radial	0.09442308	0.05139415
##	47	0.045	radial	0.09698718	0.05247373
##	48	0.050	radial	0.09955128	0.04911933
##	49	0.055	radial	0.09705128	0.05110386
##	50	0.060	radial	0.09448718	0.04698543
##	51	0.065	radial	0.09705128	0.04502436
##	52	0.070	radial	0.09705128	0.04502436
##	53	0.075	radial	0.09448718	0.04698543
##	54	0.080	radial	0.09705128	0.04502436
##	55	0.085	radial	0.09705128	0.04502436
##	56	0.090	radial	0.09961538	0.04915849
##	57	0.095	radial	0.09961538	0.04915849
##	58	0.100	radial	0.09961538	0.04915849
##	59	0.105	radial	0.09961538	0.04915849
##	60	0.110	radial	0.09961538	0.04915849
##	61	0.115	radial	0.09961538	0.04915849
##	62	0.120	radial	0.09961538	0.04915849
##	63	0.125	radial	0.10211538	0.04694056
##	64	0.130	radial	0.10467949	0.04766255
##	65	0.135	radial	0.10211538	0.04694056
##	66	0.140	radial	0.10211538	0.04694056
##	67	0.145	radial	0.10211538	0.04694056
##	68	0.150	radial	0.10211538	0.04694056
##	69	0.155	radial	0.10211538	0.04694056
##	70	0.160	radial	0.10211538	0.04694056
##	71	0.165	radial	0.10211538	0.04694056
##	72	0.170	radial	0.10467949	0.04766255
##	73	0.175	radial	0.10467949	0.04766255
##	74	0.180	radial	0.10467949	0.04766255
##	75	0.185	radial	0.10467949	0.04766255
##	76	0.190	radial	0.10467949	0.04766255
##	77	0.195	radial	0.10467949	0.04766255
##	78	0.200	radial	0.10980769	0.04219104

```

## 79 0.010 polynomial 0.23448718 0.09030846
## 80 0.015 polynomial 0.20121795 0.08974585
## 81 0.020 polynomial 0.18615385 0.08766790
## 82 0.025 polynomial 0.14801282 0.07728810
## 83 0.030 polynomial 0.13782051 0.07278833
## 84 0.035 polynomial 0.12756410 0.06833666
## 85 0.040 polynomial 0.11243590 0.07583767
## 86 0.045 polynomial 0.09448718 0.04999836
## 87 0.050 polynomial 0.09448718 0.05553599
## 88 0.055 polynomial 0.10217949 0.05934496
## 89 0.060 polynomial 0.09961538 0.05204577
## 90 0.065 polynomial 0.10474359 0.05868270
## 91 0.070 polynomial 0.10467949 0.05191916
## 92 0.075 polynomial 0.10474359 0.05326919
## 93 0.080 polynomial 0.10474359 0.05045197
## 94 0.085 polynomial 0.10730769 0.05510932
## 95 0.090 polynomial 0.10730769 0.05510932
## 96 0.095 polynomial 0.10730769 0.05510932
## 97 0.100 polynomial 0.10987179 0.05927953
## 98 0.105 polynomial 0.10730769 0.05510932
## 99 0.110 polynomial 0.10730769 0.05510932
## 100 0.115 polynomial 0.10730769 0.05510932
## 101 0.120 polynomial 0.10730769 0.05510932
## 102 0.125 polynomial 0.10730769 0.05510932
## 103 0.130 polynomial 0.10730769 0.05510932
## 104 0.135 polynomial 0.10730769 0.05510932
## 105 0.140 polynomial 0.10987179 0.05927953
## 106 0.145 polynomial 0.10987179 0.05927953
## 107 0.150 polynomial 0.11243590 0.05824145
## 108 0.155 polynomial 0.11243590 0.05824145
## 109 0.160 polynomial 0.11500000 0.06313457
## 110 0.165 polynomial 0.10987179 0.05927953
## 111 0.170 polynomial 0.11243590 0.06420732
## 112 0.175 polynomial 0.11243590 0.05824145
## 113 0.180 polynomial 0.11500000 0.06313457
## 114 0.185 polynomial 0.11500000 0.06313457
## 115 0.190 polynomial 0.11500000 0.06313457
## 116 0.195 polynomial 0.11243590 0.05824145
## 117 0.200 polynomial 0.11500000 0.05832292

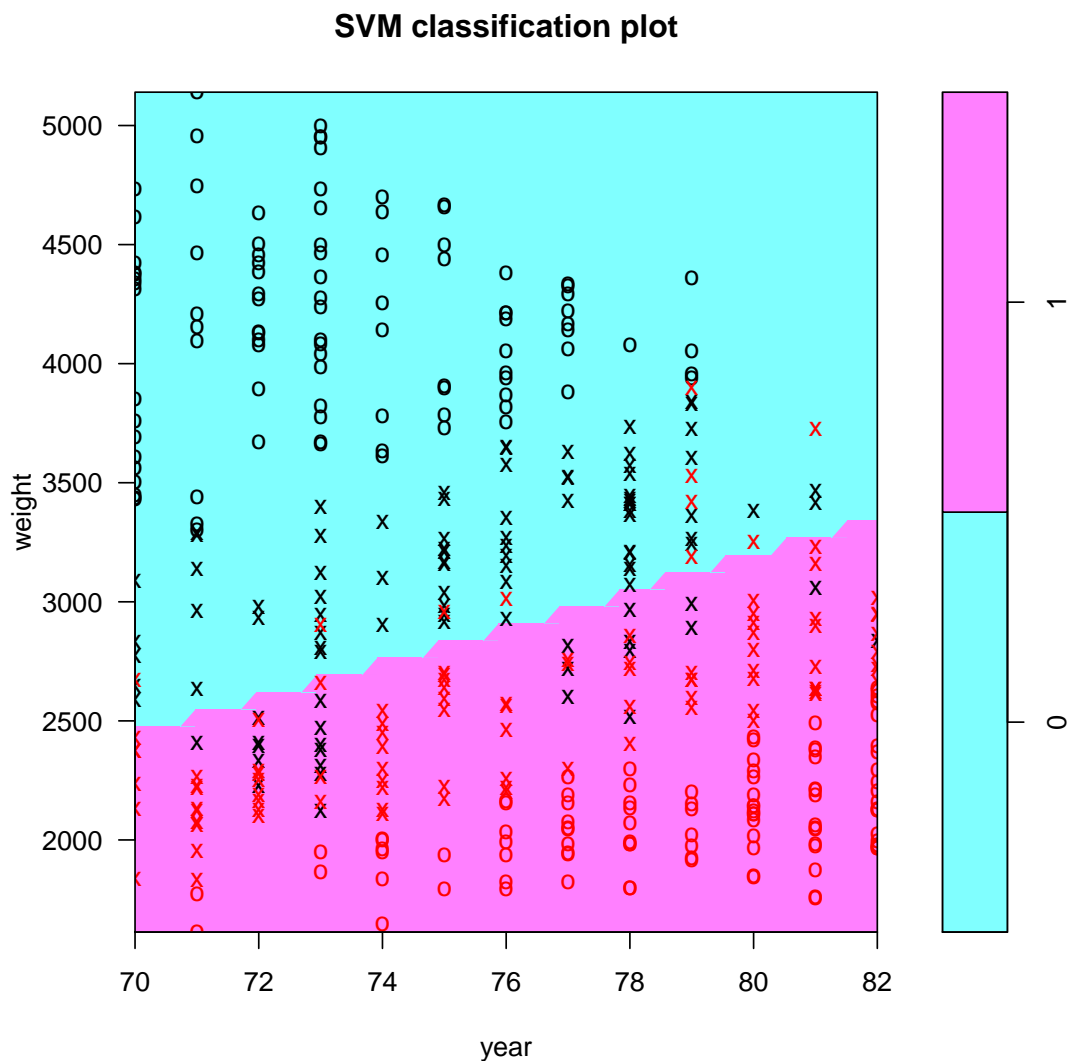
```

The best tuning parameters from the output of the 10-fold cross validation are 0.025 for cost and linear kernel, with the lowest cross-validation error 0.08423077.

- iv. Make some plots to back up your assertions in (b) and (c). When  $p > 2$ , you can use the `plot()` function to create plots displaying pairs of variables at a time. For

example, to plot horsepower and year, the syntax is `plot(svm.result, data=Auto, horsepower~year)`.

```
svm.b <- svm(mileage~weight+year, data = Auto, kernel = 'linear', cost = 0.02)
svm.c <- svm(mileage~weight+year, data = Auto, kernel = 'linear', cost = 0.025)
plot(svm.b, data=Auto, weight~year)
```



```
plot(svm.c, data=Auto, weight~year)
```

